

```
#include <iostream>
```

```
#include <vector>
```

```
#include <queue>
```

```
#include <stack>
```

```
#include <omp.h>
```

```
#include <chrono>
```

```
#include <cstdlib>
```

```
using namespace std;
```

```
using namespace std::chrono;
```

```
class Graph {
```

```
public:
```

```
    int V;
```

```
    vector<vector<int>> adj;
```

```
    Graph(int V) {
```

```
        this->V = V;
```

```
        adj.resize(V);
```

```
    }
```

```
    void addEdge(int u, int v) {
```

```
        adj[u].push_back(v);
```

```
        adj[v].push_back(u);
```

```
    }
```

```
    void generateRandomGraph(int edges);
```

```
    void sequentialBFS(int start);
```

```
    void parallelBFS(int start);
```

```
    void sequentialDFS(int start);
```

```
    void parallelDFS(int start);
```

```
};
```

```
void Graph::generateRandomGraph(int edges) {  
    srand(time(0));  
    for (int i = 0; i < edges; i++) {  
        int u = rand() % V;  
        int v = rand() % V;  
        if (u != v) {  
            addEdge(u, v);  
        }  
    }  
}
```

```
void Graph::sequentialBFS(int start) {  
    vector<bool> visited(V, false);  
    queue<int> q;  
    visited[start] = true;  
    q.push(start);  
  
    while (!q.empty()) {  
        int node = q.front();  
        q.pop();  
  
        for (int neighbor : adj[node]) {  
            if (!visited[neighbor]) {  
                visited[neighbor] = true;  
                q.push(neighbor);  
            }  
        }  
    }  
}
```

```

void Graph::parallelBFS(int start) {
    vector<bool> visited(V, false);
    queue<int> q;
    visited[start] = true;
    q.push(start);

    while (!q.empty()) {
        int size = q.size();
        vector<int> levelNodes;

        #pragma omp parallel for shared(visited, q)
        for (int i = 0; i < size; i++) {
            int node;
            #pragma omp critical
            {
                if (!q.empty()) {
                    node = q.front();
                    q.pop();
                }
            }

            for (int neighbor : adj[node]) {
                if (!visited[neighbor]) {
                    visited[neighbor] = true;
                    #pragma omp critical
                    levelNodes.push_back(neighbor);
                }
            }
        }
    }
}

```

```
    }  
    }  
}
```

```
    for (int node : levelNodes) {  
        q.push(node);  
    }  
}  
}
```

```
void Graph::sequentialDFS(int start) {  
    vector<bool> visited(V, false);  
    stack<int> s;  
    s.push(start);  
  
    while (!s.empty()) {  
        int node = s.top();  
        s.pop();  
        if (!visited[node]) {  
            visited[node] = true;  
            for (int neighbor : adj[node]) {  
                if (!visited[neighbor]) {  
                    s.push(neighbor);  
                }  
            }  
        }  
    }  
}
```

```

void Graph::parallelDFS(int start) {
    vector<bool> visited(V, false);
    stack<int> s;
    s.push(start);

    #pragma omp parallel
    {
        while (!s.empty()) {
            int node;
            #pragma omp critical
            {
                if (!s.empty()) {
                    node = s.top();
                    s.pop();
                }
            }

            if (!visited[node]) {
                visited[node] = true;
                #pragma omp parallel for
                for (int i = 0; i < adj[node].size(); i++) {
                    int neighbor = adj[node][i];
                    if (!visited[neighbor]) {
                        #pragma omp critical
                        s.push(neighbor);
                    }
                }
            }
        }
    }
}

```

```

    }
}

int main() {
    int V, E;

    cout << "Enter the number of vertices: ";
    cin >> V;

    cout << "Enter the number of edges: ";
    cin >> E;

    if (E > (V * (V - 1)) / 2) {
        cout << "Too many edges for the given number of vertices. Adjusting to maximum possible
edges.\n";
        E = (V * (V - 1)) / 2;
    }

    Graph g(V);
    g.generateRandomGraph(E);

    auto start = high_resolution_clock::now();
    g.sequentialBFS(0);
    auto stop = high_resolution_clock::now();
    auto seqBFS_time = duration_cast<microseconds>(stop - start);

    start = high_resolution_clock::now();
    g.parallelBFS(0);
    stop = high_resolution_clock::now();
    auto parBFS_time = duration_cast<microseconds>(stop - start);

```

```
start = high_resolution_clock::now();
g.sequentialDFS(0);
stop = high_resolution_clock::now();
auto seqDFS_time = duration_cast<microseconds>(stop - start);

start = high_resolution_clock::now();
g.parallelDFS(0);
stop = high_resolution_clock::now();
auto parDFS_time = duration_cast<microseconds>(stop - start);

cout << "Sequential BFS Time: " << seqBFS_time.count() << " microseconds" << endl;
cout << "Parallel BFS Time: " << parBFS_time.count() << " microseconds" << endl;
cout << "Speedup for BFS: " << (double)seqBFS_time.count() / parBFS_time.count() << endl;

cout << "Sequential DFS Time: " << seqDFS_time.count() << " microseconds" << endl;
cout << "Parallel DFS Time: " << parDFS_time.count() << " microseconds" << endl;
cout << "Speedup for DFS: " << (double)seqDFS_time.count() / parDFS_time.count() << endl;

return 0;
}
```

Output

Clear

```
Enter the number of vertices: 500
Enter the number of edges: 42000
Sequential BFS Time: 5656 microseconds
Parallel BFS Time: 5574 microseconds
Speedup for BFS: 1.01471
Sequential DFS Time: 11177 microseconds
Parallel DFS Time: 9949 microseconds
Speedup for DFS: 1.12343
```

```
=== Code Execution Successful ===
```



```

#include <iostream>

#include <chrono>

#include <omp.h>

#include <vector>

using namespace std;

using namespace std::chrono;

class Sorting {
private:
    vector<int> arr;

    int n;

    void merge(vector<int>& arr, int start, int mid, int end) {
        vector<int> left(arr.begin() + start, arr.begin() + mid + 1);
        vector<int> right(arr.begin() + mid + 1, arr.begin() + end + 1);
        int i = 0, j = 0, k = start;
        while (i < left.size() && j < right.size()) {
            if (left[i] <= right[j]) {
                arr[k++] = left[i++];
            } else {
                arr[k++] = right[j++];
            }
        }
        while (i < left.size()) arr[k++] = left[i++];
        while (j < right.size()) arr[k++] = right[j++];
    }

public:
    Sorting(vector<int> inputArr) : arr(inputArr), n(inputArr.size()) {}

    void bubbleSort() {

```

```

for (int i = 0; i < n - 1; i++) {
    for (int j = 0; j < n - i - 1; j++) {
        if (arr[j] > arr[j + 1]) {
            swap(arr[j], arr[j + 1]);
        }
    }
}
}

```

```

void mergeSort(int start, int end) {
    if (start < end) {
        int mid = start + (end - start) / 2;
        mergeSort(start, mid);
        mergeSort(mid + 1, end);
        merge(arr, start, mid, end);
    }
}

```

```

void parallelBubbleSort() {
    bool sorted = false;
    while (!sorted) {
        sorted = true;
        #pragma omp parallel for shared(arr, sorted)
        for (int i = 0; i < n - 1; i += 2) {
            if (arr[i] > arr[i + 1]) {
                swap(arr[i], arr[i + 1]);
                sorted = false;
            }
        }
        #pragma omp parallel for shared(arr, sorted)
        for (int i = 1; i < n - 1; i += 2) {

```

```

        if (arr[i] > arr[i + 1]) {
            swap(arr[i], arr[i + 1]);
            sorted = false;
        }
    }
}

```

```

void parallelMergeSort(int start, int end) {
    if (start < end) {
        int mid = start + (end - start) / 2;
        #pragma omp parallel sections
        {
            #pragma omp section
            parallelMergeSort(start, mid);
            #pragma omp section
            parallelMergeSort(mid + 1, end);
        }
        merge(arr, start, mid, end);
    }
}

```

```

void displayArray() {
    for (int i : arr) {
        cout << i << " ";
    }
    cout << endl;
}

```

```

vector<int> getArray() {
    return arr;
}

```

```

    }
};

int main() {
    cout << "Enter number of elements: ";
    int n;
    cin >> n;

    vector<int> inputArr(n);
    for (int i = 0; i < n; i++) {
        inputArr[i] = rand() % 100;
        cout << inputArr[i] << " ";
    }
    cout << "\n\n";

    Sorting sorter(inputArr);

    vector<int> originalArr = sorter.getArray();

    cout << "Sequential Execution:\n\n";

    cout << "Bubble Sort: ";
    auto start = high_resolution_clock::now();
    sorter.bubbleSort();
    auto end = high_resolution_clock::now();
    sorter.displayArray();
    double seq_bubble_time = duration<double, milli>(end - start).count();
    cout << "TIME TAKEN: " << seq_bubble_time << " ms\n";

    sorter = Sorting(originalArr);
    cout << "\nMerge Sort: ";

```

```

start = high_resolution_clock::now();
sorter.mergeSort(0, n - 1);
end = high_resolution_clock::now();
sorter.displayArray();
double seq_merge_time = duration<double, milli>(end - start).count();
cout << "TIME TAKEN: " << seq_merge_time << " ms\n";

cout << "\nParallel Execution:\n\n";

sorter = Sorting(originalArr);
cout << "Parallel Bubble Sort: ";
start = high_resolution_clock::now();
sorter.parallelBubbleSort();
end = high_resolution_clock::now();
sorter.displayArray();
double par_bubble_time = duration<double, milli>(end - start).count();
cout << "TIME TAKEN: " << par_bubble_time << " ms\n";
cout << "Speedup Factor (Bubble Sort): " << seq_bubble_time / par_bubble_time << "\n";

sorter = Sorting(originalArr);
cout << "\nParallel Merge Sort: ";
start = high_resolution_clock::now();
sorter.parallelMergeSort(0, n - 1);
end = high_resolution_clock::now();
sorter.displayArray();
double par_merge_time = duration<double, milli>(end - start).count();
cout << "TIME TAKEN: " << par_merge_time << " ms\n";
cout << "Speedup Factor (Merge Sort): " << seq_merge_time / par_merge_time << "\n";

return 0;
}

```

[illegible]

```

Sequential Execution:
Bubble Sort: 0 0 0 0 1 1 1 2 2 2 2 3 3 3 4 4 4 4 5 5 5 5 6 6 7 7 8 8 8 8 9 9 10 10 11 11 11 11 12 12 13 13 13 13 13 14 14 15 15 15 16 17 17 17 18 18
18 18 19 19 19 19 20 20 21 21 21 21 21 22 22 22 22 23 23 24 24 24 24 25 25 25 26 26 26 26 27 27 27 27 28 28 28 28 28 28 29 29 29 29
29 29 29 29 29 29 29 30 30 30 31 31 32 32 32 32 33 33 34 34 34 35 35 35 36 36 36 36 37 37 37 38 38 38 39 39 39 40 40 40 40 41 41 41
1 42 42 42 42 42 43 43 43 43 44 44 44 45 45 46 46 46 47 48 48 48 49 49 50 50 50 50 51 51 51 52 52 53 53 54 54 55 55 55 56 56
56 56 56 56 57 57 58 58 58 59 59 59 60 60 61 62 62 62 63 63 64 64 65 65 66 67 67 67 67 68 68 68 68 69 69 69 69 70 70 70
70 71 71 71 72 72 72 73 73 73 73 74 75 75 76 76 76 77 77 78 78 79 79 79 80 80 81 81 81 81 82 82 82 83 83 84 84 84 85 85 86 86
86 86 86 86 87 87 87 88 88 88 88 89 89 90 90 90 90 90 91 91 92 92 92 92 93 93 93 93 94 94 95 95 95 96 96 96 97 97 97 97 98 98 98 99
99 99 99 99 99 99
Time Taken: 2.13815 ms

```

[illegible]

TIME TAKEN: 0.615425 ms
Parallel Execution:

[illegible][illegible][illegible]

```

      speedup factor (merse sort) = 1.6439
      error estimate: 0.07100 min

```

```
#include <iostream>
```

```
#include <vector>
```

```
#include <cstdlib>
```

```
#include <ctime>
```

```
#include <chrono>
```

```
#include <omp.h>
```

```
using namespace std;
```

```
class ParallelMinMax {
```

```
private:
```

```
    vector<int> arr;
```

```
    int size;
```

```
    int min_seq, max_seq, min_par, max_par;
```

```
    long long sum_seq, sum_par;
```

```
    double avg_seq, avg_par;
```

```
    double time_min_seq, time_max_seq, time_sum_seq, time_avg_seq;
```

```
    double time_min_par, time_max_par, time_sum_par, time_avg_par;
```

```
public:
```

```
    ParallelMinMax(int size) : size(size), min_seq(0), max_seq(0), min_par(0), max_par(0), sum_seq(0),  
sum_par(0), avg_seq(0), avg_par(0),
```

```
    time_min_seq(0), time_max_seq(0), time_sum_seq(0), time_avg_seq(0), time_min_par(0),  
time_max_par(0), time_sum_par(0), time_avg_par(0) {
```

```
        arr.resize(size);
```

```
    }
```

```
    void generateRandomArray() {
```

```
        srand(time(0));
```

```
        for (int i = 0; i < size; i++) {
```

```
            arr[i] = rand() % 100 + 1;
```

```
        }
```

```
}
```

```
void computeMinSequential() {  
    auto start = chrono::high_resolution_clock::now();  
    min_seq = arr[0];  
    for (int i = 1; i < size; i++) {  
        if (arr[i] < min_seq) min_seq = arr[i];  
    }  
    auto end = chrono::high_resolution_clock::now();  
    time_min_seq = chrono::duration<double>(end - start).count();  
}
```

```
void computeMaxSequential() {  
    auto start = chrono::high_resolution_clock::now();  
    max_seq = arr[0];  
    for (int i = 1; i < size; i++) {  
        if (arr[i] > max_seq) max_seq = arr[i];  
    }  
    auto end = chrono::high_resolution_clock::now();  
    time_max_seq = chrono::duration<double>(end - start).count();  
}
```

```
void computeSumSequential() {  
    auto start = chrono::high_resolution_clock::now();  
    sum_seq = 0;  
    for (int i = 0; i < size; i++) {  
        sum_seq += arr[i];  
    }  
    auto end = chrono::high_resolution_clock::now();  
    time_sum_seq = chrono::duration<double>(end - start).count();  
}
```



```

void computeAvgSequential() {
    auto start = chrono::high_resolution_clock::now();
    avg_seq = static_cast<double>(sum_seq) / size;
    auto end = chrono::high_resolution_clock::now();
    time_avg_seq = chrono::duration<double>(end - start).count();
}

```

```

void computeMinParallel() {
    auto start = chrono::high_resolution_clock::now();
    min_par = arr[0];
    #pragma omp parallel for reduction(min:min_par)
    for (int i = 1; i < size; i++) {
        if (arr[i] < min_par) min_par = arr[i];
    }
    auto end = chrono::high_resolution_clock::now();
    time_min_par = chrono::duration<double>(end - start).count();
}

```

```

void computeMaxParallel() {
    auto start = chrono::high_resolution_clock::now();
    max_par = arr[0];
    #pragma omp parallel for reduction(max:max_par)
    for (int i = 1; i < size; i++) {
        if (arr[i] > max_par) max_par = arr[i];
    }
    auto end = chrono::high_resolution_clock::now();
    time_max_par = chrono::duration<double>(end - start).count();
}

```

```

void computeSumParallel() {

```

```
    auto start = chrono::high_resolution_clock::now();  
    sum_par = 0;  
    #pragma omp parallel for reduction(+:sum_par)  
    for (int i = 0; i < size; i++) {  
        sum_par += arr[i];  
    }  
    auto end = chrono::high_resolution_clock::now();  
    time_sum_par = chrono::duration<double>(end - start).count();  
}
```

```
void computeAvgParallel() {  
    auto start = chrono::high_resolution_clock::now();  
    avg_par = static_cast<double>(sum_par) / size;  
    auto end = chrono::high_resolution_clock::now();  
    time_avg_par = chrono::duration<double>(end - start).count();  
}
```

```
void sequentialComputation() {  
    computeMinSequential();  
    computeMaxSequential();  
    computeSumSequential();  
    computeAvgSequential();  
}
```

```
void parallelComputation() {  
    computeMinParallel();  
    computeMaxParallel();  
    computeSumParallel();  
    computeAvgParallel();  
}
```

```

void displayResults() {
    cout << "---- Sequential Computation ----\n";
    cout << "Min: " << min_seq << " | Time: " << time_min_seq << " sec\n";
    cout << "Max: " << max_seq << " | Time: " << time_max_seq << " sec\n";
    cout << "Sum: " << sum_seq << " | Time: " << time_sum_seq << " sec\n";
    cout << "Average: " << avg_seq << " | Time: " << time_avg_seq << " sec\n\n";

    cout << "---- Parallel Computation ----\n";
    cout << "Min: " << min_par << " | Time: " << time_min_par << " sec\n";
    cout << "Max: " << max_par << " | Time: " << time_max_par << " sec\n";
    cout << "Sum: " << sum_par << " | Time: " << time_sum_par << " sec\n";
    cout << "Average: " << avg_par << " | Time: " << time_avg_par << " sec\n\n";

    cout << "---- Speedup Factors ----\n";
    cout << "Speedup (Min): " << (time_min_seq / time_min_par) << "x\n";
    cout << "Speedup (Max): " << (time_max_seq / time_max_par) << "x\n";
    cout << "Speedup (Sum): " << (time_sum_seq / time_sum_par) << "x\n";
    cout << "Speedup (Average): " << (time_avg_seq / time_avg_par) << "x\n";
}
};

```

```

int main() {
    int size;
    cout << "Enter array size: ";
    cin >> size;
    if (size <= 0) {
        cout << "Invalid size!" << endl;
        return 1;
    }
}

```

```

ParallelMinMax pm(size);

```

```
pm.generateRandomArray();  
pm.sequentialComputation();  
pm.parallelComputation();  
pm.displayResults();  
  
return 0;  
}
```

Output

[Clear](#)

```
Enter array size: 1000  
---- Sequential Computation ----  
Min: 1 | Time: 4.45e-06 sec  
Max: 100 | Time: 3.94e-06 sec  
Sum: 51026 | Time: 3.83e-06 sec  
Average: 51.026 | Time: 3e-08 sec  
  
---- Parallel Computation ----  
Min: 1 | Time: 4.33e-06 sec  
Max: 100 | Time: 3.85e-06 sec  
Sum: 51026 | Time: 4.01e-06 sec  
Average: 51.026 | Time: 3e-08 sec  
  
---- Speedup Factors ----  
Speedup (Min): 1.02771x  
Speedup (Max): 1.02338x  
Speedup (Sum): 0.955112x  
Speedup (Average): 1x  
  
=== Code Execution Successful ===
```

```

#include <cuda.h>

#include <stdio.h>

#include <stdlib.h>

#include <math.h>

#include <time.h>

#include <iostream>

#include "cuda_runtime.h"

using namespace std;

#define N 1000

__global__ void vectorAdd(int *d_a, int *d_b, int *d_c, int n) {
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n) {
        d_c[i] = d_a[i] + d_b[i];
    }
}

__global__ void matrixMultiplyKernel(float *a, float *b, float *c) {
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    if (row < N && col < N) {
        float sum = 0;
        for (int i = 0; i < N; i++) {
            sum += a[row * N + i] * b[i * N + col];
        }
        c[row * N + col] = sum;
    }
}

```

```

class VectorAddition {
public:
    void performVectorAddition() {
        int *a, *b, *c, *d;
        int *d_a, *d_b, *d_c;
        size_t size = N * sizeof(int);

        a = (int *)malloc(size);
        b = (int *)malloc(size);
        c = (int *)malloc(size);
        d = (int *)malloc(size);

        srand(time(NULL));
        for (int i = 0; i < N; i++) {
            a[i] = rand() % 100;
            b[i] = rand() % 100;
        }

        clock_t start_cpu = clock();
        for (int i = 0; i < N; i++) {
            c[i] = a[i] + b[i];
        }
        clock_t end_cpu = clock();
        double cpu_time = (double)(end_cpu - start_cpu) / CLOCKS_PER_SEC;

        cudaMalloc((void **)&d_a, size);
        cudaMalloc((void **)&d_b, size);
        cudaMalloc((void **)&d_c, size);
        cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
        cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
    }
};

```

```

int threadsPerBlock = 256;

int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;


cudaEvent_t start, stop;
cudaEventCreate(&start);
cudaEventCreate(&stop);


cudaEventRecord(start);
vectorAdd<<<blocksPerGrid, threadsPerBlock>>>(d_a, d_b, d_c, N);
cudaEventRecord(stop);


cudaMemcpy(d, d_c, size, cudaMemcpyDeviceToHost);
cudaEventSynchronize(stop);

float gpu_time = 0;
cudaEventElapsedTime(&gpu_time, start, stop);


bool match = true;
for (int i = 0; i < N; i++) {
    if (c[i] != d[i]) {
        match = false;
        break;
    }
}


printf("CPU Time: %.6f s\n", cpu_time);
printf("GPU Time: %.6f ms\n", gpu_time);
printf("Speedup Factor: %.2f\n", (cpu_time) * 10000 / gpu_time);
printf("Arrays Match: %s\n", match ? "Yes" : "No");


free(a); free(b); free(c); free(d);
cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);

```

```
        cudaEventDestroy(start);  
        cudaEventDestroy(stop);  
    }  
};
```

```
class MatrixMultiplier {  
private:  
    float *hostA, *hostB, *hostC, *hostD;  
    float *devA, *devB, *devC;  
    int size;  
    float cpuTime, gpuTime;  
public:  
    MatrixMultiplier() {  
        size = N * N * sizeof(float);  
        hostA = (float *)malloc(size);  
        hostB = (float *)malloc(size);  
        hostC = (float *)malloc(size);  
        hostD = (float *)malloc(size);  
        cudaMalloc((void **)&devA, size);  
        cudaMalloc((void **)&devB, size);  
        cudaMalloc((void **)&devC, size);  
        cpuTime = gpuTime = 0.0;  
    }  
  
    ~MatrixMultiplier() {  
        free(hostA);  
        free(hostB);  
        free(hostC);  
        free(hostD);  
        cudaFree(devA);  
        cudaFree(devB);  
    }  
};
```



```
    cudaFree(devC);  
}
```

```
void initializeMatrices() {  
    for (int i = 0; i < N * N; i++) {  
        hostA[i] = rand() % 100;  
        hostB[i] = rand() % 100;  
    }  
}
```

```
void gpuMatrixMultiplication() {  
    cudaMemcpy(devA, hostA, size, cudaMemcpyHostToDevice);  
    cudaMemcpy(devB, hostB, size, cudaMemcpyHostToDevice);  
  
    dim3 dimBlock(16, 16);  
    dim3 dimGrid((N + 15) / 16, (N + 15) / 16);  
  
    clock_t tic = clock();  
    matrixMultiplyKernel<<<dimGrid, dimBlock>>>(devA, devB, devC);  
    cudaDeviceSynchronize();  
    clock_t toc = clock();  
  
    gpuTime = ((float)(toc - tic)) / CLOCKS_PER_SEC;  
    cudaMemcpy(hostC, devC, size, cudaMemcpyDeviceToHost);  
}
```

```
void cpuMatrixMultiplication() {  
    clock_t tic = clock();  
    for (int i = 0; i < N; i++) {  
        for (int j = 0; j < N; j++) {  
            float sum = 0;
```

```

for (int k = 0; k < N; k++) {
    sum += hostA[i * N + k] * hostB[k * N + j];
}
hostD[i * N + j] = sum;
}
}

clock_t toc = clock();
cpuTime = ((float)(toc - tic)) / CLOCKS_PER_SEC;
}

bool verifyEquality() {
    float tolerance = 1e-5;
    for (int i = 0; i < N * N; i++) {
        if (fabs(hostC[i] - hostD[i]) > tolerance) {
            printf("Mismatch at index %d: GPU = %f, CPU = %f\n", i, hostC[i], hostD[i]);
            return false;
        }
    }
    return true;
}

void printResults() {
    printf("CPU Time: %f seconds\n", cpuTime);
    printf("GPU Time: %f seconds\n", gpuTime);
    if (gpuTime > 0) {
        printf("Speed-Up Factor: %.2f x\n", (cpuTime) / gpuTime);
    } else {
        printf("Speed-Up Factor: N/A (GPU time too small)\n");
    }
}

```

```

    }

};

int main() {
    int choice;

    VectorAddition vectorAdder;

    MatrixMultiplier matrixMultiplier;
    matrixMultiplier.initializeMatrices();
    matrixMultiplier.cpuMatrixMultiplication();
    matrixMultiplier.gpuMatrixMultiplication();
    bool success = matrixMultiplier.verifyEquality();

    do {
        cout << "\nMenu:" << endl;
        cout << "1. Vector Addition" << endl;
        cout << "2. Matirx Multiplication" << endl;
        cout << "3. Exit" << endl;
        cout << "Enter your choice: ";
        cin >> choice;
        switch (choice) {
            case 1:
                vectorAdder.performVectorAddition();
                break;
            case 2:

                matrixMultiplier.printResults();
                printf("Verification: %s\n", success ? "true" : "false");
                break;
            case 3:

```

```

    cout << "Exiting..." << endl;

    break;

    default:

    cout << "Invalid choice!" << endl;

    }

    } while (choice != 3);

    return 0;

}

```

```

(base) dypcoe-student@admin1-MS-7D48:~/Downloads$ nvcc cuda_program.cu
(base) dypcoe-student@admin1-MS-7D48:~/Downloads$ ./a.out
Menu:
1. Vector Addition
2. Matrix Multiplication
3. Exit
Enter your choice: 1
Enter the number of elements (N): 100000
Vector Addition Results:
CPU Time: 0.000993 s
GPU Time: 0.037472 ms
Speedup Factor: 26.50
Arrays Match: Yes

Menu:
1. Vector Addition
2. Matrix Multiplication
3. Exit
Enter your choice: 2
Enter the size of the matrix (N x N): 25
Matrix Multiplication Results:
CPU Time: 0.000142 seconds
GPU Time: 0.000054 seconds
Speed-Up Factor: 2.03 x
Verification: true

Menu:
1. Vector Addition
2. Matrix Multiplication
3. Exit
Enter your choice: 3
Exiting program...
(base) dypcoe-student@admin1-MS-7D48:~/Downloads$

```

```

#include <iostream>

#include <vector>

#include <cmath>

#include <chrono>

#include <cuda.h>

#define N 1000 // number of data points

// CUDA Kernel for parallel sum calculation
__global__ void computeSums(const float* x, const float* y, float* sumX, float* sumY, float* sumXY,
float* sumX2, int n) {

    int i = blockIdx.x * blockDim.x + threadIdx.x;

    if (i < n) {

        atomicAdd(sumX, x[i]);

        atomicAdd(sumY, y[i]);

        atomicAdd(sumXY, x[i] * y[i]);

        atomicAdd(sumX2, x[i] * x[i]);

    }

}

// CPU implementation of linear regression
void linearRegressionCPU(const float* x, const float* y, int n, float& b0, float& b1) {

    float sumX = 0, sumY = 0, sumXY = 0, sumX2 = 0;

    for (int i = 0; i < n; i++) {

        sumX += x[i];

        sumY += y[i];

        sumXY += x[i] * y[i];

        sumX2 += x[i] * x[i];

    }

    b1 = (n * sumXY - sumX * sumY) / (n * sumX2 - sumX * sumX);

    b0 = (sumY - b1 * sumX) / n;

```

```
}
```

```
int main() {  
    // Host data  
    std::vector<float> h_x(N), h_y(N);  
    float b0_cpu, b1_cpu;  
    float b0_gpu, b1_gpu;  
  
    // Generate synthetic data:  $y = 2.5 + 1.2x$   
    for (int i = 0; i < N; ++i) {  
        h_x[i] = i;  
        h_y[i] = 2.5f + 1.2f * i;  
    }  
  
    // === CPU Linear Regression ===  
    auto start_cpu = std::chrono::high_resolution_clock::now();  
    linearRegressionCPU(h_x.data(), h_y.data(), N, b0_cpu, b1_cpu);  
    auto end_cpu = std::chrono::high_resolution_clock::now();  
    float cpu_time = std::chrono::duration<float, std::milli>(end_cpu - start_cpu).count();  
  
    // === GPU Linear Regression ===  
    float *d_x, *d_y, *d_sumX, *d_sumY, *d_sumXY, *d_sumX2;  
    cudaMalloc(&d_x, N * sizeof(float));  
    cudaMalloc(&d_y, N * sizeof(float));  
    cudaMalloc(&d_sumX, sizeof(float));  
    cudaMalloc(&d_sumY, sizeof(float));  
    cudaMalloc(&d_sumXY, sizeof(float));  
    cudaMalloc(&d_sumX2, sizeof(float));  
  
    // Initialize device memory for sums  
    cudaMemset(d_sumX, 0, sizeof(float));
```

```

cudaMemset(d_sumY, 0, sizeof(float));
cudaMemset(d_sumXY, 0, sizeof(float));
cudaMemset(d_sumX2, 0, sizeof(float));

cudaMemcpy(d_x, h_x.data(), N * sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(d_y, h_y.data(), N * sizeof(float), cudaMemcpyHostToDevice);

int threadsPerBlock = 256;
int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;

cudaEvent_t start_gpu, stop_gpu;
cudaEventCreate(&start_gpu);
cudaEventCreate(&stop_gpu);
cudaEventRecord(start_gpu);

computeSums<<<blocksPerGrid, threadsPerBlock>>>(d_x, d_y, d_sumX, d_sumY, d_sumXY, d_sumX2,
N);

cudaEventRecord(stop_gpu);
cudaEventSynchronize(stop_gpu);

float gpu_time = 0;
cudaEventElapsedTime(&gpu_time, start_gpu, stop_gpu);

float sumX, sumY, sumXY, sumX2;
cudaMemcpy(&sumX, d_sumX, sizeof(float), cudaMemcpyDeviceToHost);
cudaMemcpy(&sumY, d_sumY, sizeof(float), cudaMemcpyDeviceToHost);
cudaMemcpy(&sumXY, d_sumXY, sizeof(float), cudaMemcpyDeviceToHost);
cudaMemcpy(&sumX2, d_sumX2, sizeof(float), cudaMemcpyDeviceToHost);

b1_gpu = (N * sumXY - sumX * sumY) / (N * sumX2 - sumX * sumX);
b0_gpu = (sumY - b1_gpu * sumX) / N;

bool verified = fabs(b0_cpu - b0_gpu) < 1e-2 && fabs(b1_cpu - b1_gpu) < 1e-2;

```

```

std::cout << "=== Linear Regression using CUDA ===\n";

std::cout << "CPU Time: " << cpu_time << " ms\n";

std::cout << "GPU Time: " << gpu_time << " ms\n";

std::cout << "Speedup Factor: " << cpu_time / gpu_time << "x\n";

std::cout << "Verification: " << (verified ? "PASSED ✓" : "FAILED ✗") << "\n";

std::cout << "Equation (CPU): y = " << b0_cpu << " + " << b1_cpu << " * x\n";

std::cout << "Equation (GPU): y = " << b0_gpu << " + " << b1_gpu << " * x\n";

cudaFree(d_x);

cudaFree(d_y);

cudaFree(d_sumX);

cudaFree(d_sumY);

cudaFree(d_sumXY);

cudaFree(d_sumX2);

cudaEventDestroy(start_gpu);

cudaEventDestroy(stop_gpu);

return 0;

}

```

The screenshot shows a terminal window with the following output:

```

(base) dypcoe-student@admin1-M5-7D48:~/Downloads$ nvcc linear_reg.cu
(base) dypcoe-student@admin1-M5-7D48:~/Downloads$ ./a.out
=== Linear Regression using CUDA ===
CPU Time: 0.017419 ms
GPU Time: 0.019104 ms
Speedup Factor: 0.911799x
Verification: PASSED ✓
Equation (CPU): y = 2.40887 + 1.2 * x
Equation (GPU): y = 2.49738 + 1.20001 * x
(base) dypcoe-student@admin1-M5-7D48:~/Downloads$

```

The terminal window is titled "dypcoe-student@admin1-M5-7D48: ~/Downloads". The desktop background is dark purple. On the left, there is a sidebar with icons for Activities, Terminal, and various file managers. On the right, there is a file manager window showing a list of files and folders, including "Cloud Computing 3.pdf", "MP handwriting note Un...", "MP Unit-1 St-II.pdf", "student.py", "1.py", "archive(1).zip", "Comp March 2024 Ne...", "Comp March 2024 Ne...", "Comp March 2024 Ne...", "lp-1.zip", "Screenshot from 2024-07-10 15...", and "Screenshot from 2024-07-10 15...".